

```
In [1]: import pandas as pd
import numpy as np

import matplotlib as pyplot
import matplotlib.pyplot as plt; plt.rcParams()
import matplotlib.pyplot as plt
%matplotlib inline

import seaborn as sns
```

```
In [2]: df = pd.read_csv('train.csv')
```

## Review the data set

```
In [3]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 550068 entries, 0 to 550067
Data columns (total 12 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   User_ID                               550068 non-null  int64
1   Product_ID                           550068 non-null  object
2   Gender                               550068 non-null  object
3   Age                                   550068 non-null  object
4   Occupation                           550068 non-null  int64
5   City_Category                       550068 non-null  object
6   Stay_In_Current_City_Years          550068 non-null  object
7   Marital_Status                      550068 non-null  int64
8   Product_Category_1                  550068 non-null  int64
9   Product_Category_2                  376430 non-null  float64
10  Product_Category_3                  166821 non-null  float64
11  Purchase                             550068 non-null  int64
dtypes: float64(2), int64(5), object(5)
memory usage: 50.4+ MB
```

```
In [4]: df.describe()
```

```
Out[4]:
```

	User_ID	Occupation	Marital_Status	Product_Category_1	Product_Category_2	Product_Categor
<b>count</b>	5.500680e+05	550068.000000	550068.000000	550068.000000	376430.000000	166821.000
<b>mean</b>	1.003029e+06	8.076707	0.409653	5.404270	9.842329	12.668
<b>std</b>	1.727592e+03	6.522660	0.491770	3.936211	5.086590	4.125
<b>min</b>	1.000001e+06	0.000000	0.000000	1.000000	2.000000	3.000
<b>25%</b>	1.001516e+06	2.000000	0.000000	1.000000	5.000000	9.000
<b>50%</b>	1.003077e+06	7.000000	0.000000	5.000000	9.000000	14.000
<b>75%</b>	1.004478e+06	14.000000	1.000000	8.000000	15.000000	16.000
<b>max</b>	1.006040e+06	20.000000	1.000000	20.000000	18.000000	18.000

```
In [5]: df.head()
```

```
Out[5]:
```


	User_ID	Product_ID	Gender	Age	Occupation	City_Category	Stay_In_Current_City_Years	Marital_Status
0	1000001	P00069042	F	0-17	10	A	2	0
1	1000001	P00248942	F	0-17	10	A	2	0
2	1000001	P00087842	F	0-17	10	A	2	0
3	1000001	P00085442	F	0-17	10	A	2	0
4	1000002	P00285442	M	55+	16	C	4+	0

```
In [6]: df['Product_Category_1'].unique()
```

```
Out[6]: array([ 3,  1, 12,  8,  5,  4,  2,  6, 14, 11, 13, 15,  7, 16, 18, 10, 17,
           9, 20, 19], dtype=int64)
```

```
In [7]: sns.heatmap(df.isnull(), cbar=False)
```

```
Out[7]: <matplotlib.axes._subplots.AxesSubplot at 0x1f9bebf3108>
```




```
In [8]: # Clear null values
df.isnull().sum()
```

```
Out[8]: User_ID          0
Product_ID          0
Gender              0
Age                0
Occupation          0
City_Category       0
Stay_In_Current_City_Years  0
Marital_Status      0
Product_Category_1   0
Product_Category_2  173638
Product_Category_3  383247
Purchase            0
dtype: int64
```


```
In [9]: sns.countplot(df['Age'], hue=df['Gender'])
```

```
Out[9]: <matplotlib.axes._subplots.AxesSubplot at 0x1f9c10e1c08>
```




```
In [10]: sns.countplot(df['Stay_In_Current_City_Years'])
```

```
Out[10]: <matplotlib.axes._subplots.AxesSubplot at 0x1f9c16b47c8>
```




```
In [11]: sns.countplot(df['Occupation'])
```

```
Out[11]: <matplotlib.axes._subplots.AxesSubplot at 0x1f9c1488ec8>
```




```
In [12]: sns.countplot(df['Product_Category_1'])
```

```
Out[12]: <matplotlib.axes._subplots.AxesSubplot at 0x1f9c1d63288>
```



```
In [13]: sns.distplot(df['Purchase'])
```

```
Out[13]: <matplotlib.axes._subplots.AxesSubplot at 0x1f9c13d6288>
```



We have labeled data - can use supervised learning We want to predict how much \$\$\$ people spend - regression task

When we think about sales - consider demographic (age/gender/income/home), store (location/size/ads/product stock)

```
In [14]: print(df['Product_ID'].unique())
print(len(df['Product_ID'].unique()))

['P00069042' 'P00248942' 'P00087842' ... 'P00370293' 'P00371644'
 'P00370853']
3631
```

## Prepare data

```
In [15]: X = df.drop(['Purchase', 'Product_Category_2', 'Product_Category_3'], axis=1)
```

```
In [16]: for cat in ['Age', 'City_Category', 'Stay_In_Current_City_Years']:
        tmp = pd.get_dummies(X[cat], prefix=cat, drop_first=True)
        X.drop(cat, axis=1, inplace=True)
        X = pd.concat([X, tmp], axis=1)
```

```
In [17]: X['Gender'] = pd.get_dummies(X['Gender'], drop_first=True)
```

```
In [18]: from sklearn.preprocessing import LabelEncoder

        for c in ['Product_ID']:
            enc = LabelEncoder()
            X[c] = enc.fit_transform(X[c])
```

```
In [19]: X["Occupation"] = X["Occupation"].astype('category')
```

```
In [20]: X.head()
```

```
Out[20]:
```

	User_ID	Product_ID	Gender	Occupation	Marital_Status	Product_Category_1	Age_18-25	Age_26-35	Age_36-45
0	1000001	672	0	10	0	3	0	0	0
1	1000001	2376	0	10	0	1	0	0	0
2	1000001	852	0	10	0	12	0	0	0
3	1000001	828	0	10	0	12	0	0	0
4	1000002	2734	1	16	0	8	0	0	0

```
In [21]: X.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 550068 entries, 0 to 550067
Data columns (total 18 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   User_ID                               550068 non-null  int64
1   Product_ID                           550068 non-null  int32
2   Gender                               550068 non-null  uint8
3   Occupation                           550068 non-null  category
4   Marital_Status                       550068 non-null  int64
5   Product_Category_1                   550068 non-null  int64
6   Age_18-25                           550068 non-null  uint8
7   Age_26-35                           550068 non-null  uint8
8   Age_36-45                           550068 non-null  uint8
9   Age_46-50                           550068 non-null  uint8
10  Age_51-55                           550068 non-null  uint8
11  Age_55+                             550068 non-null  uint8
12  City_Category_B                     550068 non-null  uint8
13  City_Category_C                     550068 non-null  uint8
14  Stay_In_Current_City_Years_1        550068 non-null  uint8
15  Stay_In_Current_City_Years_2        550068 non-null  uint8
16  Stay_In_Current_City_Years_3        550068 non-null  uint8
17  Stay_In_Current_City_Years_4+       550068 non-null  uint8
dtypes: category(1), int32(1), int64(3), uint8(13)
memory usage: 22.0 MB
```

```
In [22]: y = df["Purchase"]
```

```
In [23]: X = X.drop(["User_ID"], axis=1)
```

```
In [24]: from sklearn.preprocessing import StandardScaler
```

```
In [25]: scaler = StandardScaler()
Xs = scaler.fit_transform(X)
```

## Regression Model Comparison

```
In [26]: import sklearn
import xgboost as xgb

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor
from sklearn.ensemble import GradientBoostingRegressor
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsRegressor
from sklearn.linear_model import Ridge
from sklearn.linear_model import Lasso
from sklearn.preprocessing import PolynomialFeatures
from sklearn import svm
```

```
In [27]: from sklearn import metrics

def get_metrics(reg, y_test, y_pred):
    res = []

    res.append(str(reg.__class__.__name__))
    mae = metrics.mean_absolute_error(y_test, y_pred)
    res.append(mae)
    mse = metrics.mean_squared_error(y_test, y_pred)
    res.append(mse)
    rmse = np.sqrt(metrics.mean_squared_error(y_test, y_pred))
    res.append(rmse)
    r2 = metrics.r2_score(y_test, y_pred)
    res.append(r2)

    return res

def show_metrics(res):
    print(res[0])
    print(' MAE : ', res[1])
    print(' MSE : ', res[2])
    print(' RMSE: ', res[3])
    print(' R^2 : ', res[4])
```

```
In [28]: X_train, X_test, y_train, y_test = train_test_split(Xs, y, test_size = 1/3, random_
state = 0)
```

```
In [29]: # Start with polynomial regression
pf = PolynomialFeatures(degree=2)
X_poly_train = pf.fit_transform(X_train)
X_poly_test = pf.fit_transform(X_test)
```

```
In [30]: lr = LinearRegression()
lr.fit(X_poly_train, y_train)
```

```
Out[30]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)
```

```
In [31]: y_pred = lr.predict(X_poly_test)
print(y_pred)
```

```
[13555.  5193.  8640. ...  7041.  7711. 12155.]
```

```
In [32]: results = []
m = get_metrics(lr, y_test, y_pred)
m[0] = 'PolynomialRegression'
show_metrics(m)
results.append(m)
```

```
PolynomialRegression
MAE : 3309.958206985318
MSE : 19261757.748472918
RMSE: 4388.821908949248
R^2 : 0.23849121551232466
```

```
In [33]: regressors = [
    Ridge(alpha=.5),
    Lasso(alpha=0.1),
    LinearRegression(),
    DecisionTreeRegressor(),
    RandomForestRegressor(),
    GradientBoostingRegressor(),
    xgb.XGBRegressor(objective='reg:squarederror', colsample_bytree = 0.3, learning_rate = 0.1,
                      max_depth = 5, alpha = 10, n_estimators = 10),
    # svm.SVC(kernel='rbf', gamma=0.7, C=1.0),
    KNeighborsRegressor(n_neighbors=3)
]
```

```
### Commented out regressors are processing over 100 minutes
```

```
In [34]: for r in regressors:
          print('=====')
          print('=====' )
          print('Running %s' % (r))
          r.fit(X_train, y_train)
          # print('Accuracy score for %s is %s' % (r, r.score(X_test, y_test) * 100))
          m = get_metrics(r, y_test, r.predict(X_test))
          show_metrics(m)
          results.append(m)
          print('=====')
          print('')
```



```

=====
Running Ridge(alpha=0.5, copy_X=True, fit_intercept=True, max_iter=None,
              normalize=False, random_state=None, solver='auto', tol=0.001)
Ridge
MAE : 3593.7867622984763
MSE : 21952039.018481284
RMSE: 4685.300312518002
R^2 : 0.1321316170475142
=====

Running Lasso(alpha=0.1, copy_X=True, fit_intercept=True, max_iter=1000,
              normalize=False, positive=False, precompute=False, random_state=None,
              selection='cyclic', tol=0.0001, warm_start=False)
Lasso
MAE : 3593.7848869242807
MSE : 21952043.356074445
RMSE: 4685.300775411803
R^2 : 0.1321314455618513
=====

Running LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize
=False)
LinearRegression
MAE : 3593.786431714285
MSE : 21952039.006256834
RMSE: 4685.300311213448
R^2 : 0.13213161753080482
=====

Running DecisionTreeRegressor(ccp_alpha=0.0, criterion='mse', max_depth=None,
                              max_features=None, max_leaf_nodes=None,
                              min_impurity_decrease=0.0, min_impurity_split=None,
                              min_samples_leaf=1, min_samples_split=2,
                              min_weight_fraction_leaf=0.0, presort='deprecated',
                              random_state=None, splitter='best')
DecisionTreeRegressor
MAE : 2630.8304546378868
MSE : 13872403.005075026
RMSE: 3724.5674923506253
R^2 : 0.45155801000791973
=====

Running RandomForestRegressor(bootstrap=True, ccp_alpha=0.0, criterion='mse',
                              max_depth=None, max_features='auto', max_leaf_nodes=None,
                              max_samples=None, min_impurity_decrease=0.0,
                              min_impurity_split=None, min_samples_leaf=1,
                              min_samples_split=2, min_weight_fraction_leaf=0.0,
                              n_estimators=100, n_jobs=None, oob_score=False,
                              random_state=None, verbose=0, warm_start=False)
RandomForestRegressor
MAE : 2104.1033457206127

```

## Try ANN

```
In [35]: from sklearn.pipeline import Pipeline
        from sklearn.preprocessing import MinMaxScaler
        from keras.optimizers import Adam
        from keras.models import Sequential
        from keras.layers import Dense
        from keras.wrappers.scikit_learn import KerasRegressor
```

Using TensorFlow backend.

```
In [36]: model = Sequential()
        model.add(Dense(17, input_dim=17, kernel_initializer='normal', activation='relu'))
        model.add(Dense(27, activation='relu'))
        model.add(Dense(11, activation='relu'))
        model.add(Dense(3, activation='relu'))
        model.add(Dense(1, activation='relu'))

        optimizer = Adam(lr=0.35)
        model.compile(loss='mse', optimizer=optimizer, metrics=['mse', 'mae'])
```

```
In [37]: model.fit(X_train, y_train, epochs=300, batch_size=5000, verbose=1)
```

Epoch 1/300  
366712/366712 [=====] - 0s 1us/step - loss: 26250183.59  
21 - mse: 26250192.0000 - mae: 3795.1655

Epoch 2/300  
366712/366712 [=====] - 0s 1us/step - loss: 18074555.00  
54 - mse: 18074556.0000 - mae: 3146.5696

Epoch 3/300  
366712/366712 [=====] - 0s 1us/step - loss: 17228205.35  
37 - mse: 17228208.0000 - mae: 3051.3696

Epoch 4/300  
366712/366712 [=====] - 0s 1us/step - loss: 16850510.21  
17 - mse: 16850510.0000 - mae: 3011.1182

Epoch 5/300  
366712/366712 [=====] - 0s 1us/step - loss: 17581072.35  
68 - mse: 17581074.0000 - mae: 3091.6904

Epoch 6/300  
366712/366712 [=====] - 0s 1us/step - loss: 16342514.04  
16 - mse: 16342516.0000 - mae: 2964.5261

Epoch 7/300  
366712/366712 [=====] - 0s 1us/step - loss: 15924482.97  
80 - mse: 15924482.0000 - mae: 2927.9519

Epoch 8/300  
366712/366712 [=====] - 0s 1us/step - loss: 15766800.70  
47 - mse: 15766799.0000 - mae: 2912.5049

Epoch 9/300  
366712/366712 [=====] - 0s 1us/step - loss: 16444965.80  
77 - mse: 16444965.0000 - mae: 2979.8274

Epoch 10/300  
366712/366712 [=====] - 0s 1us/step - loss: 16069738.93  
82 - mse: 16069736.0000 - mae: 2937.9956

Epoch 11/300  
366712/366712 [=====] - 0s 1us/step - loss: 15010027.29  
08 - mse: 15010027.0000 - mae: 2852.6082

Epoch 12/300  
366712/366712 [=====] - 0s 1us/step - loss: 12847770.65  
33 - mse: 12847771.0000 - mae: 2656.2136

Epoch 13/300  
366712/366712 [=====] - 0s 1us/step - loss: 12201562.65  
25 - mse: 12201563.0000 - mae: 2587.5259

Epoch 14/300  
366712/366712 [=====] - 0s 1us/step - loss: 11538222.25  
72 - mse: 11538222.0000 - mae: 2544.5303

Epoch 15/300  
366712/366712 [=====] - 0s 1us/step - loss: 11169928.65  
88 - mse: 11169929.0000 - mae: 2525.0774

Epoch 16/300  
366712/366712 [=====] - 0s 1us/step - loss: 10341835.55  
42 - mse: 10341836.0000 - mae: 2428.5405

Epoch 17/300  
366712/366712 [=====] - 0s 1us/step - loss: 10086132.49  
00 - mse: 10086132.0000 - mae: 2391.1746

Epoch 18/300  
366712/366712 [=====] - 0s 1us/step - loss: 10143938.56  
89 - mse: 10143938.0000 - mae: 2388.6416

Epoch 19/300  
366712/366712 [=====] - 0s 1us/step - loss: 10246179.95  
99 - mse: 10246182.0000 - mae: 2402.2083

Epoch 20/300  
366712/366712 [=====] - 0s 1us/step - loss: 9908738.943  
6 - mse: 9908739.0000 - mae: 2369.1973

Epoch 21/300  
366712/366712 [=====] - 0s 1us/step - loss: 9922185.573  
1 - mse: 9922186.0000 - mae: 2365.0410

Epoch 22/300

Out[37]: <keras.callbacks.callbacks.History at 0x1fa184128c8>


```
In [38]: m = get_metrics(model, y_test, model.predict(X_test))
m[0] = 'ANN'
show_metrics(m)
results.append(m)
```

```
ANN
MAE : 2280.8142783018657
MSE : 9094922.777238183
RMSE: 3015.778966906922
R^2 : 0.640434498266234
```

```
In [39]: res = pd.DataFrame(np.array(results))
res.columns = ['Name', 'MAE', 'MSE', 'RMSE', 'R^2']
```

```
In [40]: sns.barplot(data=res, y='Name', x='R^2')
```

Out[40]: <matplotlib.axes.\_subplots.AxesSubplot at 0x1fa24883208>



```
In [41]: df_r = pd.DataFrame(results)
df_r.head()
```

```
Out[41]:
```


		0	1	2	3	4
0	PolynomialRegression	3309.958207	1.926176e+07	4388.821909	0.238491	
1	Ridge	3593.786762	2.195204e+07	4685.300313	0.132132	
2	Lasso	3593.784887	2.195204e+07	4685.300775	0.132131	
3	LinearRegression	3593.786432	2.195204e+07	4685.300311	0.132132	
4	DecisionTreeRegressor	2630.830455	1.387240e+07	3724.567492	0.451558	

## Clustering Algorithm Comparison

### K-Means

```
In [42]: from sklearn.cluster import KMeans
wcss = []
for i in range(1, 11):
    print("Running KMeans with %s clusters" % (i))
    kmeans = KMeans(n_clusters = i, init = 'k-means++', random_state = 42)
    kmeans.fit(X)
    wcss.append(kmeans.inertia_)
plt.plot(range(1, 11), wcss)
plt.title('The Elbow Method')
plt.xlabel('Number of clusters')
plt.ylabel('WCSS') # sum of squares of the distances of each data point in all clusters to their centroids
plt.show()
```

```
Running KMeans with 1 clusters
Running KMeans with 2 clusters
Running KMeans with 3 clusters
Running KMeans with 4 clusters
Running KMeans with 5 clusters
Running KMeans with 6 clusters
Running KMeans with 7 clusters
Running KMeans with 8 clusters
Running KMeans with 9 clusters
Running KMeans with 10 clusters
```



## PCA

```
In [43]: from sklearn.decomposition import PCA
pca = PCA(n_components = 3)
X_pca = pca.fit_transform(X)
explained_variance = pca.explained_variance_ratio_
```

```
In [44]: print(explained_variance) # % of variance explained by each of the selected components
[9.99941231e-01 4.15276886e-05 1.50365374e-05]
```

```
In [45]: print(pca.components_)
```

```
[[ 9.99999957e-01  6.78015193e-06  4.78419549e-05  5.64613034e-06
  2.90559777e-04 -5.80060188e-06 -7.43338167e-06  4.90866276e-06
  2.69169851e-06  2.62798511e-06  2.23825144e-06 -2.70419324e-06
  3.99272802e-06 -2.55125162e-07  1.24345339e-06  1.95922614e-07
 -1.46872062e-06]
 [-4.57358766e-05  7.79265580e-03  9.99906755e-01  1.83354708e-03
 -7.73317533e-03 -5.71348649e-03 -1.68022493e-03  3.58707265e-03
  8.55606948e-04  1.20291330e-03  1.30603065e-03 -8.90990100e-04
  2.46657997e-03 -1.58035596e-04 -7.60300602e-04  1.51493781e-03
  7.71763687e-04]
 [-2.90943163e-04 -5.05095526e-03  7.73703507e-03  2.48351949e-03
  9.99938267e-01 -3.49261764e-03 -2.25884654e-03  1.17554421e-03
  1.80871965e-03  1.67135689e-03  1.67204132e-03 -7.35715933e-04
 -1.03593996e-03  7.69793887e-04 -7.45522006e-04 -1.04656767e-03
  6.59269406e-04]]
```

## Feature Selection

```
In [46]: from sklearn.feature_selection import VarianceThreshold
Xvt = X.copy(deep=True)
sel = VarianceThreshold(threshold=(.8 * (1 - .8)))
sel.fit_transform(Xvt)
```

```
Out[46]: array([[672, 0, 10, ..., 0, 0, 0],
 [2376, 0, 10, ..., 0, 0, 0],
 [852, 0, 10, ..., 0, 0, 0],
 ...,
 [3568, 0, 15, ..., 1, 0, 0],
 [3568, 0, 1, ..., 0, 1, 0],
 [3566, 0, 0, ..., 1, 0, 0]], dtype=object)
```

```

In [47]: regressors = [
    LinearRegression(),
    RandomForestRegressor(),
    KNeighborsRegressor(n_neighbors=3)
]

X_train, X_test, y_train, y_test = train_test_split(Xvt, y, test_size = 1/3, random
_state = 0)

for r in regressors:
    print('=====')
    print('Running %s' % (r))
    r.fit(X_train, y_train)

    m = get_metrics(r, y_test, r.predict(X_test))
    show_metrics(m)
    print('=====')
    print('')

=====

Running LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize
=False)
LinearRegression
MAE : 3593.7864317142844
MSE : 21952039.006256834
RMSE: 4685.300311213448
R^2 : 0.13213161753080482
=====

Running RandomForestRegressor(bootstrap=True, ccp_alpha=0.0, criterion='mse',
max_depth=None, max_features='auto', max_leaf_nodes=None,
max_samples=None, min_impurity_decrease=0.0,
min_impurity_split=None, min_samples_leaf=1,
min_samples_split=2, min_weight_fraction_leaf=0.0,
n_estimators=100, n_jobs=None, oob_score=False,
random_state=None, verbose=0, warm_start=False)
RandomForestRegressor
MAE : 2103.764583410758
MSE : 8209169.894973392
RMSE: 2865.1648983912587
R^2 : 0.6754525173659389
=====

Running KNeighborsRegressor(algorithm='auto', leaf_size=30, metric='minkowski',
metric_params=None, n_jobs=None, n_neighbors=3, p=2,
weights='uniform')
KNeighborsRegressor
MAE : 2253.43161572751
MSE : 9501839.293711567
RMSE: 3082.505359883672
R^2 : 0.6243471553614997
=====

```



**RFE**

```

In [48]: from sklearn.feature_selection import RFE

regressors = [
    LinearRegression(),
    # RandomForestRegressor(),
    DecisionTreeRegressor(),
    GradientBoostingRegressor()
]

X_train, X_test, y_train, y_test = train_test_split(Xvt, y, test_size = 1/3, random
_state = 0)

for r in regressors:
    print('=====')
    print('Running %s' % (r))
    rfe = RFE(r, 3)
    ref = rfe.fit(X_train, y_train)
    print("Regressor: " + str(r.__class__.__name__))
    print("Support")
    print(rfe.support_)
    print("Ranking")
    print(rfe.ranking_)
    print('=====')
    print('')

```

```

=====
=====
Running LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize
=False)
Regressor: LinearRegression
Support
[False True False False True False False False False False False False
 True False False False False]
Ranking
[15 1 14 9 1 7 6 4 5 2 3 8 1 13 10 12 11]
=====
=====

=====
=====
Running DecisionTreeRegressor(ccp_alpha=0.0, criterion='mse', max_depth=None,
                             max_features=None, max_leaf_nodes=None,
                             min_impurity_decrease=0.0, min_impurity_split=None,
                             min_samples_leaf=1, min_samples_split=2,
                             min_weight_fraction_leaf=0.0, presort='deprecated',
                             random_state=None, splitter='best')
Regressor: DecisionTreeRegressor
Support
[ True False True False True False False False False False False False
 False False False False False]
Ranking
[ 1 7 1 4 1 11 5 9 13 14 15 3 12 2 8 6 10]
=====
=====

=====
=====
Running GradientBoostingRegressor(alpha=0.9, ccp_alpha=0.0, criterion='friedman_
mse',
                                init=None, learning_rate=0.1, loss='ls', max_depth=3,
                                max_features=None, max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, n_estimators=100,
                                n_iter_no_change=None, presort='deprecated',
                                random_state=None, subsample=1.0, tol=0.0001,
                                validation_fraction=0.1, verbose=0, warm_start=False)
Regressor: GradientBoostingRegressor
Support
[ True False False False True False False False False False False False
 True False False False False]
Ranking
[ 1 7 2 14 1 8 9 6 13 3 5 4 1 10 11 12 15]
=====
=====

```